

## Meets Specifications

Great submission all around. You took the time to think through the design and architecture of the networks and the output performed very well. Congrats!

### Files Submitted



The submission includes all required files.

### Step 1: Detect Humans



The submission returns the percentage of the first 100 images in the dog and human face datasets with a detected human face.



The submission opines whether Haar cascades for face detection are an appropriate technique for human detection.

### Step 2: Detect Dogs



The submission returns the percentage of the first 100 images in the dog and human face datasets with a detected dog.

### Step 3: Create a CNN to Classify Dog Breeds (from Scratch)



The submission specifies a CNN architecture.

Good description of the layers used and rationale behind using them.

I've seen good performance with an additional Conv layer, larger initial kernel size, or experimenting with Dropout layers after the max pooling layers.

Two links that helped me understand how some of the parameters work.

<https://adeshpande3.github.io/adeshpande3.github.io/A-Beginner's-Guide-To-Understanding-Convolutional-Neural-Networks/>



The submission specifies the number of epochs used to train the algorithm.



The trained model attains at least 1% accuracy on the test set.

Test accuracy: 7.0574%

## Step 5: Create a CNN to Classify Dog Breeds



The submission downloads the bottleneck features corresponding to one of the Keras pre-trained models (VGG-19, ResNet-50, Inception, or Xception).

Good idea to try them all.



The submission specifies a model architecture.



The submission details why the chosen architecture succeeded in the classification task and why earlier attempts were not as successful.



The submission compiles the architecture by specifying the loss function and optimizer.



The submission uses model checkpointing to train the model and saves the model weights with the best validation loss.



The submission loads the model weights that attained the least validation loss.



Accuracy on the test set is 60% or greater.

Resnet50 81.5789% 80.9809%  
InceptionV3 79.7847% 81.2201%  
Xception 84.4498% 84.6890%



The submission includes a function that takes a file path to an image as input and returns the dog breed that is predicted by the CNN.

## Step 6: Write Your Algorithm



The submission uses the CNN from Step 5 to detect dog breed. The submission has different output for each detected image type (dog, human, other) and provides either predicted actual (or resembling) dog breed.

Forgot an output for when neither are detected. Just prints breed with no other feedback.

## Step 7: Test Your Algorithm



The submission tests at least 6 images, including at least two human and two dog images.

Good set of sample images some humans and dogs and a couple to try to mess with the algorithm. Includes good ideas for next direction to take this for improvement.